

In [1]:

```

import sys
import numpy as np
import time
import os
import logging
from collections import Counter
from datetime import datetime
import math

from tqdm import tqdm
import pickle

import torch
import torch.nn as nn
import torch.nn.functional as F
from torch import nn, optim

```

## Assignment 3: Dependency Parsing

### Estimated Time: ~10 hours

This assignment will build a neural dependency parser using PyTorch. In part 1, we will review two general neural network techniques (Adam optimization and Dropout). In part 2, we will implement and train a dependency parser using techniques from part 1.

## Part 1. Adam Optimization and Dropout

### a) Adam

Recall the SGD update rule:

$$\theta = \theta - \alpha \nabla_{\theta} J_{\text{minibatch}}(\theta)$$

where  $\theta$  is a vector containing all of the model parameters,  $J$  is the loss function,  $\nabla_{\theta} J_{\text{minibatch}}(\theta)$  is the gradient of the loss function, and  $\alpha$  is the learning rate. Adam is another possible update rule with two additional steps.

- (2 pts) First, Adam uses a trick called momentum by keep track of  $\mathbf{m}$ , a rolling average of the gradients:

$$\begin{aligned} \mathbf{m} &= \beta_1 \mathbf{m} + (1 - \beta_1) \nabla_{\theta} J_{\text{minibatch}}(\theta) \\ \theta &= \theta - \alpha \mathbf{m} \end{aligned}$$

where  $\beta_1$  is a hyperparameter between 0 and 1 (often set to 0.9). Briefly explain in 2-4 sentences (just give an intuition) how using  $\mathbf{m}$  stops the updates from varying as much and why this low variance may be helpful to learning, overall.

**Write your answer here.**

**Answer :**

- Momentum simply means that some fraction of the previous update is added to the current update and makes the current gradient not just dependent on its mini-batch gradient like SGD did.
- The momentum which is the exponential moving average of the gradient(m), is updated where the hyper-parameters control the exponential decay rates of these moving averages.
- In some cases, if the learning rate is too large, the optimized solution(diverge) will be missed. If the learning rate is too small, it will converge very small.
- The using momentum can make the learning rate more stable (not change too fast) and faster convergence.
- Moreover, high variance can have a tendency to be overly complex and cause overfitting cases but low variance can reduce the number of samples needed to obtain accurate estimates.
- (2 pts) Adam extends the idea of momentum with the trick of adaptive learning rates by keep track of  $\mathbf{v}$ , a rolling average of the magnitudes of the gradients:

$$\begin{aligned}\mathbf{m} &= \beta_1 \mathbf{m} + (1 - \beta_1) \nabla_{\theta} J_{\text{minibatch}}(\theta) \\ \mathbf{v} &= \beta_2 \mathbf{v} + (1 - \beta_2) (\nabla_{\theta} J_{\text{minibatch}}(\theta) \circ \nabla_{\theta} J_{\text{minibatch}}(\theta)) \\ \theta &= \theta - \alpha \mathbf{m} / \sqrt{\mathbf{v}}\end{aligned}$$

where  $\circ$  and  $/$  denote elementwise multiplication and division (not dot product!).  $\beta_2$  is a hyperparameter between 0 and 1 (often set to 0.99). Since Adam divides the update by  $\sqrt{\mathbf{v}}$ , what kinds of weights will receive larger update and smaller update? Give some simple example of how. Why might this help with learning?

**Write your answer here.**

**Answer :**

- The model parameters which receive small or infrequent updates will get larger updates. In the opposite situation, the model parameters which receive larger updates will have their effective learning rate reduced. Thus, we can regard adaptive learning rates as normalization of the parameter update step by element wise.
- Parameters that previous gradients are small and not volatile get larger updates. This helps model to handle with sparse gradients (merits of AdaGrad) and also non-stationary objectives (merits of RMSProp)

## b) Dropout

Dropout is a regularization technique. During training, dropout randomly sets units in the hidden layer  $\mathbf{h}$  to zero with probability  $p_{\text{drop}}$  (dropping different units each minibatch), and then multiplies  $\mathbf{h}$  by a constant  $\gamma$ . We can write this as:

$$\mathbf{h}_{\text{drop}} = \gamma \mathbf{d} \circ \mathbf{h}$$

where  $\mathbf{d} \in \{0, 1\}^{D_h}$  ( $D_h$  is the size of  $\mathbf{h}$ ) is a mask vector where each entry is 0 with probability  $p_{\text{drop}}$  and 1 with probability  $(1 - p_{\text{drop}})$ . For the gamma constant,  $\gamma$  is chosen such that the expected value of  $\mathbf{h}_{\text{drop}}$  is  $\mathbf{h}$

$$\mathbb{E}_{p_{\text{drop}}}[\mathbf{h}_{\text{drop}}]_i = h_i$$

for all  $i \in \{1, \dots, D_h\}$

- (2 pts) What must  $\gamma$  equal in term of  $p_{\text{drop}}$ ? Briefly justify your answers or show your math derivation using the equations given above.

**Write your answer here.**

**Answer :**

$$\gamma \text{ must be } \frac{1}{(1 - p_{\text{drop}})}$$

$$\mathbb{E}_{p_{\text{drop}}}[\mathbf{h}_{\text{drop}}]_i = \mathbf{h}_i$$

$$\mathbb{E}_{p_{\text{drop}}}[\gamma \mathbf{d} \circ \mathbf{h}]_i = \mathbf{h}_i$$

$$\mathbb{E}_{p_{\text{drop}}}[\mathbf{d}_i] \gamma \mathbf{h}_i = \mathbf{h}_i$$

$$\gamma = \frac{1}{\mathbb{E}_{p_{\text{drop}}}[\mathbf{d}_i]}$$

Since  $\mathbf{d} \in \{0, 1\}^{D_h}$  is a mask vector where each entry is 0 with probability  $p_{\text{drop}}$  and 1 with probability  $(1 - p_{\text{drop}})$ ,

$$\gamma = \frac{1}{(0 \cdot p_{\text{drop}} + 1(1 - p_{\text{drop}}))}$$

$$\gamma = \frac{1}{(1 - p_{\text{drop}})}$$

- (2pts) Why should dropout be applied only during training? Why should dropout NOT be applied during evaluation?

**Write your answer here.**

**Answer :**

Drop out is one of the regularizations which aim to prevent overfitting. Since overfitting only happens during training, should dropout be applied only during training. Besides, evaluation is made from train, to model being improved. But if we use drop out also at evaluation, it does re-trained with same circumstances, so that can't do well in test.

## Part 2. Neural Transition-Based Dependency Parsing

We will be implementing a neural dependency parser with the goal of maximizing the performance on the UAS (Unlabeled Attachment Score) metric.

A dependency parser analyzes the grammatical structure of a sentence, establishing relationships between head words, and words which modify those heads. There are multiple types of dependency parsers, including transition-based parsers, graph-based parsers, and feature-based parsers. Your implementation will be a transition-based parser, which incrementally builds up a parse one step at a time. At every step, it maintains a partial parse, which is represented as follows:

- A **stack** of words that are currently being processed
- A **buffer** of words yet to be processed.
- A **list of dependencies** predicted by the parser

Initially, the stack only contains ROOT, the dependencies list is empty, and the buffer contains all words of the sentence in order. At each step, the parser applies a transition to the partial parse until its buffer is empty and the stack size is 1. The following transitions can be applied:

- **SHIFT**: removes the first word from the buffer and pushes it onto the stack.

- LEFTARC: marks the second (second msot recently aded) item on the stack as a dependent of the first item and removes the second item from the stack, adding a *first\_word* → *second\_word* dependency to the dependency list.
- RIGHTARC: marks the first (second msot recently aded) item on the stack as a dependent of the second item and removes the first item from the stack, adding a *second\_word* → *first\_word* dependency to the dependency list.

On each step, your parser will decide among the three transitions using a neural network classifier.

- (4 pts) Go through the sequence of transitions needed for parsing the sentence *I parsed this sentence correctly*. The dependency tree for the sentence is shown below. At each step, give the configuration of the stack and buffer, as well as what transition was applied this step and what new dependency was added (if any). The first three steps are provided below as an example.

Stack	Buffer	New dependency	Transition
[ROOT]	[I, parsed, this, sentence, correctly]		Init
[ROOT, I]	[parsed, this, sentence, correctly]		SHIFT
[ROOT, I, parsed]	[this, sentence, correctly]		SHIFT
[ROOT, parsed]	[this, sentence, correctly]	parsed → I	LEFTARC

**Write your answer here.**

**Answer :**

Stack	Buffer	New dependency	Transition
[ROOT]	[I, parsed, this, sentence, correctly]		Init
[ROOT, I]	[parsed, this, sentence, correctly]		SHIFT
[ROOT, I, parsed]	[this, sentence, correctly]		SHIFT
[ROOT, parsed]	[this, sentence, correctly]	parsed → I	LEFTARC
[ROOT, parsed, this]	[sentence, correctly]		SHIFT
[ROOT, parsed, this, sentence]	[correctly]		SHIFT
[ROOT, parsed, sentence]	[correctly]	sentence → this	LEFTARC
[ROOT, parsed]	[correctly]	parsed → sentence	RIGHTARC
[ROOT, parsed, correctly]			SHIFT
[ROOT, parsed]		parsed → correctly	RIGHTARC
[ROOT]		ROOT → parsed	RIGHTARC

- (2 pts) A sentence containing *n* words will be parsed in how many steps (in terms of *n*)? Briefly explain in 1-2 sentences why.

**Write your answer here.**

**Answer :**

A sentence containing *n* words will be parsed in  $2n$  steps. Because we will eventually move all words form the Buffer to Stack (the SHIFT) step and them remove all words from the Stack(the ARC steps).

## Part 3: Parser

- (6pts) Implement the `__init__` and `parse_step` functions in the `PartialParse` class below. This implements the transition mechanics your parser will use.

Test your function by running `test_parse_step()` followed by `test_parse()`

- (8pts) Our network will predict which transition should be applied next to a partial parse. We could use it to parse a single sentence by applying predicted transitions until the parse is complete. However, neural networks run much more efficiently when making predictions about batches of data (i.e., predicting the next transition for any different partial parses simultaneously). We can parse sentences in minibatches with the following algorithm

**Input:** `sentences` , a list of sentences to be parsed and `model` , our model that makes parse decisions

1. Initialize `partial_pares` as a list of `PartialPares` , one for each sentence in `sentences`
2. Initailize `unfinished_pares` as a shallow copy of `partial_pares`
3. **while** `unfinished_pares` is not empty **do**
  - Take the first `batch_size` parses in `unfinished_pares` as a minibatch
  - Use the `model` to predict the next transition for each partial parse in the minitbatch
  - Peform a parse step on each partial parse in the minibatch with its predicted transition
  - Remove the completed (empty buffer and stack of size 1) parses from `unfinished_pares`

**Return:** The dependencies for each parse in `partial_pares`

Implement this algorithm in the `minibatch_parse` function below.

Test your function by running `test_minibatch_parse()`

In [2]:

```

class PartialParse(object):
    def __init__(self, sentence):
        """Initializes this partial parse.
        @param sentence (list of str): The sentence to be parsed as a list of words.
            Your code should not modify the sentence.

        """
        # The sentence being parsed is kept for bookkeeping purposes. Do not alter it.
        self.sentence = sentence

        # YOUR CODE HERE (3 Lines)
        # Your code should initialize the following fields:
        # self.stack: The current stack represented as a list with the top of the stack at the
        # last element of the list.
        # self.buffer: The current buffer represented as a list with the first item at the
        # buffer as the first item of the list
        # self.dependencies: The list of dependencies produced so far. Represented as a list of
        # tuples where each tuple is of the form (head, dependent).
        # Order for this list doesn't matter.
        ###
        # Note: The root token should be represented with the string "ROOT"
        ###

        self.stack = ["ROOT"]
        self.buffer = self.sentence.copy()
        self.dependencies = []

        # END YOUR CODE

    def parse_step(self, transition):
        """Performs a single parse step by applying the given transition to
        this partial parse.
        @param transition (str): A string that equals "S", "LA", or "RA"
            representing the shift, left-arc, and right-arc
            transitions. You can assume the provided
            transition is a legal transition.

        """
        # YOUR CODE HERE (~7-10 Lines)
        # TODO:
        # Implement a single parsing step, i.e. the logic for the following as
        # described in the pdf handout:
        # 1. Shift
        # 2. Left Arc
        # 3. Right Arc

        if transition == "S":
            self.stack += [self.buffer.pop(0)]
        else:
            last = self.stack[-1]
            prev = self.stack[-2]
            if transition == "RA":
                self.dependencies += [(prev, last)]
                self.stack.pop(-1)
            if transition == "LA":
                self.dependencies += [(last, prev)]
                self.stack.pop(-2)

```

```

# END YOUR CODE

def parse(self, transitions):
    """Applies the provided transitions to this PartialParse
    @param transitions (list of str): The list of transitions in the order
        they should be applied.
    @return dependencies (list of string tuples): The list of dependencies
        produced when parsing the
        sentence. Represented as
        a list of tuples where each
        tuple is of the form
        (head, dependent).

    """
    for transition in transitions:
        self.parse_step(transition)
    return self.dependencies

def is_completed(self):
    return (len(self.buffer) == 0) and (len(self.stack) == 1)

def minibatch_parse(sentences, model, batch_size):
    """Parses a list of sentences in minibatches using a model.
    @param sentences (list of list of str): A list of sentences to be parsed
        (each sentence is a list of words
        and each word is of type string)
    @param model (ParserModel): The model that makes parsing decisions. It is
        assumed to have a function
        model.predict(partial_parses) that takes in a
        list of PartialParses as input and
        returns a list of transitions predicted for each
        parse. That is, after calling
        transitions = model.predict(partial_parses)
        transitions[i] will be the next transition to
        apply to partial_parses[i].
    @param batch_size (int): The number of PartialParses to include in each minibatch
    @return dependencies (list of dependency lists): A list where each element
        is the dependencies list for
        parsed sentence. Ordering
        should be the same as in
        sentences (i.e.,
        dependencies[i] should
        contain the parse for
        sentences[i]).

    """
    dependencies = []

    # YOUR CODE HERE (~8-10 Lines)
    # TODO:
    # Implement the minibatch parse algorithm
    ###
    # Note: A shallow copy can be made with the "=" sign
    # in python, e.g. unfinished_parses = partial_parses[:].
    # Here `unfinished_parses` is a shallow copy of `partial_parses`.
    # In Python, a shallow copied list like `unfinished_parses` does not contain
    # new instances of the object stored in `partial_parses`. Rather both lists
    # refer to the same objects.

```

```
# In our case, `partial_parsers` contains a list of partial parses.
# `unfinished_parsers` contains references to the same objects. Thus, you
# should NOT use the `del` operator to remove objects from the
# `unfinished_parsers` list. This will free the underlying memory that
# is being accessed by `partial_parsers` and may cause your code to crash.
```

```
partial_parsers = [PartialParse(sentence) for sentence in sentences]
unfinished_parsers = partial_parsers[: ]

while unfinished_parsers:
    minibatch = unfinished_parsers[:batch_size]
    transitions = model.predict(minibatch)

    for partial_parse, transition in zip(minibatch, transitions):
        partial_parse.parse_step(transition)

    unfinished_parsers = [partial_parse for partial_parse in partial_parsers
                          if partial_parse.buffer or len(partial_parse.stack) !=
                          batch_size]

dependencies = [partial_parse.dependencies for partial_parse in partial_parsers]

# END YOUR CODE

return dependencies
```

```
def test_step(name, transition, stack, buf, deps,
              ex_stack, ex_buf, ex_deps):
    """Tests that a single parse step returns the expected output"""
    pp = PartialParse([])
    pp.stack, pp.buffer, pp.dependencies = stack, buf, deps

    pp.parse_step(transition)
    stack, buf, deps = (tuple(pp.stack), tuple(pp.buffer),
                       tuple(sorted(pp.dependencies)))
    assert stack == ex_stack, \
        "{:} test resulted in stack {:}, expected {:}".format(
            name, stack, ex_stack)
    assert buf == ex_buf, \
        "{:} test resulted in buffer {:}, expected {:}".format(
            name, buf, ex_buf)
    assert deps == ex_deps, \
        "{:} test resulted in dependency list {:}, expected {:}".format(
            name, deps, ex_deps)
    print("{:} test passed!".format(name))
```

```
def test_parse_step():
    """Simple tests for the PartialParse.parse_step function
    Warning: these are not exhaustive
    """
    test_step("SHIFT", "S", ["ROOT", "the"], ["cat", "sat"], [],
              ("ROOT", "the", "cat"), ("sat",), ())
    test_step("LEFT-ARC", "LA", ["ROOT", "the", "cat"], ["sat"], [],
              ("ROOT", "cat",), ("sat",), (("cat", "the"),))
    test_step("RIGHT-ARC", "RA", ["ROOT", "run", "fast"], [], [],
              ("ROOT", "run",), (), (("run", "fast"),))
```



```

def test_parse():
    """Simple tests for the PartialParse.parse function
    Warning: these are not exhaustive
    """
    sentence = ["parse", "this", "sentence"]
    dependencies = PartialParse(sentence).parse(
        ["S", "S", "S", "LA", "RA", "RA"])
    dependencies = tuple(sorted(dependencies))
    expected = (('ROOT', 'parse'), ('parse', 'sentence'), ('sentence', 'this'))
    assert dependencies == expected, \
        "parse test resulted in dependencies {:}, expected {:}".format(
            dependencies, expected)
    assert tuple(sentence) == ("parse", "this", "sentence"), \
        "parse test failed: the input sentence should not be modified"
    print("parse test passed!")

class DummyModel(object):
    """Dummy model for testing the minibatch_parse function
    First shifts everything onto the stack and then does exclusively right arcs if the
    the sentence is "right", "left" if otherwise.
    """

    def predict(self, partial_parses):
        return [("RA" if pp.stack[1] == "right" else "LA") if len(pp.buffer) == 0 else "no"
                for pp in partial_parses]

def test_dependencies(name, deps, ex_deps):
    """Tests the provided dependencies match the expected dependencies"""
    deps = tuple(sorted(deps))
    assert deps == ex_deps, \
        "{:} test resulted in dependency list {:}, expected {:}".format(
            name, deps, ex_deps)

def test_minibatch_parse():
    """Simple tests for the minibatch_parse function
    Warning: these are not exhaustive
    """
    sentences = [
        ["right", "arcs", "only"],
        ["right", "arcs", "only", "again"],
        ["left", "arcs", "only"],
        ["left", "arcs", "only", "again"]]
    deps = minibatch_parse(sentences, DummyModel(), 2)
    test_dependencies("minibatch_parse", deps[0],
        (('ROOT', 'right'), ('arcs', 'only'), ('right', 'arcs')))
    test_dependencies("minibatch_parse", deps[1],
        (('ROOT', 'right'), ('arcs', 'only'), ('only', 'again'), ('right', 'arcs')))
    test_dependencies("minibatch_parse", deps[2],
        (('only', 'ROOT'), ('only', 'arcs'), ('only', 'left')))
    test_dependencies("minibatch_parse", deps[3],
        (('again', 'ROOT'), ('again', 'arcs'), ('again', 'left'), ('again', 'arcs')))
    print("minibatch_parse test passed!")

```

Test your function by running `test_parse_step()` followed by `test_parse()`

In [3]:

```
#testing your parse_step
#turn on when you are ready
test_parse_step()
test_parse
```

```
SHIFT test passed!
LEFT-ARC test passed!
RIGHT-ARC test passed!
```

Out[3]:

```
<function __main__.test_parse(>
```

Test your function by running `test_minibatch_parse()`

In [4]:

```
#testing your minibatch_parse
#turn on when you are ready
test_minibatch_parse()
```

```
minibatch_parse test passed!
```

## Part 4: Neural Network

Let's train a neural network to predict, given the state of the stack, buffer, and dependencies, which transition should be applied next.

First, the model extracts a feature vector representing the current state. We will be using the feature set presented in *A Fast and Accurate Dependency Parser using Neural Networks* (Chen and Manning 2014). The function extracting these features are implemented for you here below.

This feature vector consists of a list of tokens (e.g., the last word in the stack, first word in the buffer, dependent of the second-to-last word in the stack if there is one, etc.). They can be represented as a list of integers  $\mathbf{w} = [w_1, w_2, \dots, w_m]$  where  $m$  is the number of features and each  $0 \leq w_i \leq |V|$  is the index of a token in the vocabulary ( $|V|$  is the vocabulary size). Then our network looks up an embedding for each word and concatenates them into a single input vector:

$$\mathbf{x} = [\mathbf{E}_{w_1}, \dots, \mathbf{E}_{w_m}] \in \mathbb{R}^{dm}$$

where  $\mathbf{E} \in \mathbb{R}^{|V| \times d}$  is an embedding matrix with each row  $\mathbf{E}_w$  as the vector for a particular word  $w$

We then compute our prediction as:

$$\begin{aligned}\mathbf{h} &= \text{ReLU}(\mathbf{xW} + \mathbf{b}_1) \\ \mathbf{l} &= \mathbf{hU} + \mathbf{b}_2 \\ \hat{\mathbf{y}} &= \text{softmax}(\mathbf{l})\end{aligned}$$

where  $\mathbf{h}$  is referred to as the hidden layer,  $\mathbf{l}$  is the logits,  $\hat{\mathbf{y}}$  is the predictions, and  $\text{ReLU}(z) = \max(z, 0)$ . We will then train the model to minimize cross-entropy (CE) loss:

$$J(\theta) = \text{CE}(\mathbf{y}, \hat{\mathbf{y}}) = - \sum_{i=1}^3 y_i \log \hat{y}_i$$

To compute the loss for the training set, we average this  $J(\theta)$  across all training examples. We will use UAS (Unlabeled Attachment Score) as main metric, which is computed as the ratio between number of correctly predicted dependencies and the number of total dependencies despite of the relations (our model doesn't predict this).

Below this code, you will find a skeleton code to implement this network using PyTorch. Complete the `__init__`, `embedding_lookup` and `forward` functions to implement the model. Then complete the `train_for_epoch` and `train` functions to actually train the model.

In [5]:

*#do not change this code; this is provided for you which helps extract features from*

```

P_PREFIX = '<p>:'
L_PREFIX = '<l>:'
UNK = '<UNK>'
NULL = '<NULL>'
ROOT = '<ROOT>'

def get_minibatches(data, minibatch_size, shuffle=True):
    """
    Iterates through the provided data one minibatch at a time. You can use this function to
    iterate through data in minibatches as follows:

        for inputs_minibatch in get_minibatches(inputs, minibatch_size):
            ...

    Or with multiple data sources:

        for inputs_minibatch, labels_minibatch in get_minibatches([inputs, labels],
            ...

    Args:
        data: there are two possible values:
            - a list or numpy array
            - a list where each element is either a list or numpy array
        minibatch_size: the maximum number of items in a minibatch
        shuffle: whether to randomize the order of returned data

    Returns:
        minibatches: the return value depends on data:
            - If data is a list/array it yields the next minibatch of data.
            - If data is a list of lists/arrays it returns the next minibatch of each element
              list. This can be used to iterate through multiple data sources
              (e.g., features and labels) at the same time.

    """
    list_data = type(data) is list and (type(data[0]) is list or type(data[0]) is np.ndarray)
    data_size = len(data[0]) if list_data else len(data)
    indices = np.arange(data_size)
    if shuffle:
        np.random.shuffle(indices)
    for minibatch_start in np.arange(0, data_size, minibatch_size):
        minibatch_indices = indices[minibatch_start:minibatch_start + minibatch_size]
        yield [_minibatch(d, minibatch_indices) for d in data] if list_data \
            else _minibatch(data, minibatch_indices)

def _minibatch(data, minibatch_idx):
    return data[minibatch_idx] if type(data) is np.ndarray else [data[i] for i in minibatch_idx]

def test_all_close(name, actual, expected):
    if actual.shape != expected.shape:
        raise ValueError("{} failed, expected output to have shape {} but has shape {}".format(name, expected.shape, actual.shape))
    if np.amax(np.fabs(actual - expected)) > 1e-6:
        raise ValueError("{} failed, expected {} but value is {}".format(name, expected, actual))
    else:
        print(name, "passed!")

```

```

class Config(object):
    language = 'english'
    with_punct = True
    unlabeled = True
    lowercase = True
    use_pos = True
    use_dep = True
    use_dep = use_dep and (not unlabeled)
    data_path = './data-a3'
    train_file = 'train.conll'
    dev_file = 'dev.conll'
    test_file = 'test.conll'
    embedding_file = './data-a3/en-cw.txt'

class Parser(object):
    """Contains everything needed for transition-based dependency parsing except for

def __init__(self, dataset):
    root_labels = list([l for ex in dataset
                        for (h, l) in zip(ex['head'], ex['label']) if h == 0])
    counter = Counter(root_labels)
    if len(counter) > 1:
        logging.info('Warning: more than one root label')
        logging.info(counter)
    self.root_label = counter.most_common()[0][0]
    deprel = [self.root_label] + list(set([w for ex in dataset
                                         for w in ex['label']
                                         if w != self.root_label]))
    tok2id = {L_PREFIX + l: i for (i, l) in enumerate(deprel)}
    tok2id[L_PREFIX + NULL] = self.L_NULL = len(tok2id)

    config = Config()
    self.unlabeled = config.unlabeled
    self.with_punct = config.with_punct
    self.use_pos = config.use_pos
    self.use_dep = config.use_dep
    self.language = config.language

    if self.unlabeled:
        trans = ['L', 'R', 'S']
        self.n_deprel = 1
    else:
        trans = ['L-' + l for l in deprel] + ['R-' + l for l in deprel] + ['S']
        self.n_deprel = len(deprel)

    self.n_trans = len(trans)
    self.tran2id = {t: i for (i, t) in enumerate(trans)}
    self.id2tran = {i: t for (i, t) in enumerate(trans)}

    # logging.info('Build dictionary for part-of-speech tags.')
    tok2id.update(build_dict([P_PREFIX + w for ex in dataset for w in ex['pos']],
                             offset=len(tok2id)))
    tok2id[P_PREFIX + UNK] = self.P_UNK = len(tok2id)
    tok2id[P_PREFIX + NULL] = self.P_NULL = len(tok2id)
    tok2id[P_PREFIX + ROOT] = self.P_ROOT = len(tok2id)

    # logging.info('Build dictionary for words.')
    tok2id.update(build_dict([w for ex in dataset for w in ex['word']],
                             offset=len(tok2id)))

```

```

tok2id[UNK] = self.UNK = len(tok2id)
tok2id[NULL] = self.NULL = len(tok2id)
tok2id[ROOT] = self.ROOT = len(tok2id)

self.tok2id = tok2id
self.id2tok = {v: k for (k, v) in tok2id.items()}

self.n_features = 18 + (18 if config.use_pos else 0) + (12 if config.use_dep else 0)
self.n_tokens = len(tok2id)

def vectorize(self, examples):
    vec_examples = []
    for ex in examples:
        word = [self.ROOT] + [self.tok2id[w] if w in self.tok2id
                               else self.UNK for w in ex['word']]
        pos = [self.P_ROOT] + [self.tok2id[P_PREFIX + w] if P_PREFIX + w in self.tok2id
                                else self.P_UNK for w in ex['pos']]
        head = [-1] + ex['head']
        label = [-1] + [self.tok2id[L_PREFIX + w] if L_PREFIX + w in self.tok2id
                        else -1 for w in ex['label']]
        vec_examples.append({'word': word, 'pos': pos,
                            'head': head, 'label': label})

    return vec_examples

def extract_features(self, stack, buf, arcs, ex):
    if stack[0] == "ROOT":
        stack[0] = 0

    def get_lc(k):
        return sorted([arc[1] for arc in arcs if arc[0] == k and arc[1] < k])

    def get_rc(k):
        return sorted([arc[1] for arc in arcs if arc[0] == k and arc[1] > k],
                       reverse=True)

    p_features = []
    l_features = []
    features = [self.NULL] * (3 - len(stack)) + [ex['word'][x] for x in stack[-3:]]
    features += [ex['word'][x] for x in buf[:3]] + [self.NULL] * (3 - len(buf))
    if self.use_pos:
        p_features = [self.P_NULL] * (3 - len(stack)) + [ex['pos'][x] for x in stack[-3:]]
        p_features += [ex['pos'][x] for x in buf[:3]] + [self.P_NULL] * (3 - len(buf))

    for i in range(2):
        if i < len(stack):
            k = stack[-i-1]
            lc = get_lc(k)
            rc = get_rc(k)
            llc = get_lc(lc[0]) if len(lc) > 0 else []
            rrc = get_rc(rc[0]) if len(rc) > 0 else []

            features.append(ex['word'][lc[0]] if len(lc) > 0 else self.NULL)
            features.append(ex['word'][rc[0]] if len(rc) > 0 else self.NULL)
            features.append(ex['word'][lc[1]] if len(lc) > 1 else self.NULL)
            features.append(ex['word'][rc[1]] if len(rc) > 1 else self.NULL)
            features.append(ex['word'][llc[0]] if len(llc) > 0 else self.NULL)
            features.append(ex['word'][rrc[0]] if len(rrc) > 0 else self.NULL)

            if self.use_pos:
                p_features.append(ex['pos'][lc[0]] if len(lc) > 0 else self.P_NULL)
                p_features.append(ex['pos'][rc[0]] if len(rc) > 0 else self.P_NULL)

```

```

        p_features.append(ex['pos'][lc[1]] if len(lc) > 1 else self.P_NULL)
        p_features.append(ex['pos'][rc[1]] if len(rc) > 1 else self.P_NULL)
        p_features.append(ex['pos'][llc[0]] if len(llc) > 0 else self.P_NULL)
        p_features.append(ex['pos'][rrc[0]] if len(rrc) > 0 else self.P_NULL)

    if self.use_dep:
        l_features.append(ex['label'][lc[0]] if len(lc) > 0 else self.L_NULL)
        l_features.append(ex['label'][rc[0]] if len(rc) > 0 else self.L_NULL)
        l_features.append(ex['label'][lc[1]] if len(lc) > 1 else self.L_NULL)
        l_features.append(ex['label'][rc[1]] if len(rc) > 1 else self.L_NULL)
        l_features.append(ex['label'][llc[0]] if len(llc) > 0 else self.L_NULL)
        l_features.append(ex['label'][rrc[0]] if len(rrc) > 0 else self.L_NULL)
    else:
        features += [self.NULL] * 6
        if self.use_pos:
            p_features += [self.P_NULL] * 6
        if self.use_dep:
            l_features += [self.L_NULL] * 6

    features += p_features + l_features
    assert len(features) == self.n_features
    return features

def get_oracle(self, stack, buf, ex):
    if len(stack) < 2:
        return self.n_trans - 1

    i0 = stack[-1]
    i1 = stack[-2]
    h0 = ex['head'][i0]
    h1 = ex['head'][i1]
    l0 = ex['label'][i0]
    l1 = ex['label'][i1]

    if self.unlabeled:
        if (i1 > 0) and (h1 == i0):
            return 0
        elif (i1 >= 0) and (h0 == i1) and \
            (not any([x for x in buf if ex['head'][x] == i0])):
            return 1
        else:
            return None if len(buf) == 0 else 2
    else:
        if (i1 > 0) and (h1 == i0):
            return l1 if (l1 >= 0) and (l1 < self.n_deprel) else None
        elif (i1 >= 0) and (h0 == i1) and \
            (not any([x for x in buf if ex['head'][x] == i0])):
            return l0 + self.n_deprel if (l0 >= 0) and (l0 < self.n_deprel) else None
        else:
            return None if len(buf) == 0 else self.n_trans - 1

def create_instances(self, examples):
    all_instances = []
    succ = 0
    for id, ex in enumerate(examples):
        n_words = len(ex['word']) - 1

        # arcs = {(h, t, label)}
        stack = [0]
        buf = [i + 1 for i in range(n_words)]
        arcs = []

```

```

instances = []
for i in range(n_words * 2):
    gold_t = self.get_oracle(stack, buf, ex)
    if gold_t is None:
        break
    legal_labels = self.legal_labels(stack, buf)
    assert legal_labels[gold_t] == 1
    instances.append((self.extract_features(stack, buf, arcs, ex),
                     legal_labels, gold_t))
    if gold_t == self.n_trans - 1:
        stack.append(buf[0])
        buf = buf[1:]
    elif gold_t < self.n_deprel:
        arcs.append((stack[-1], stack[-2], gold_t))
        stack = stack[:-2] + [stack[-1]]
    else:
        arcs.append((stack[-2], stack[-1], gold_t - self.n_deprel))
        stack = stack[:-1]
else:
    succ += 1
    all_instances += instances

return all_instances

def legal_labels(self, stack, buf):
    labels = ([1] if len(stack) > 2 else [0]) * self.n_deprel
    labels += ([1] if len(stack) >= 2 else [0]) * self.n_deprel
    labels += [1] if len(buf) > 0 else [0]
    return labels

def parse(self, dataset, eval_batch_size=5000):
    sentences = []
    sentence_id_to_idx = {}
    for i, example in enumerate(dataset):
        n_words = len(example['word']) - 1
        sentence = [j + 1 for j in range(n_words)]
        sentences.append(sentence)
        sentence_id_to_idx[id(sentence)] = i

    model = ModelWrapper(self, dataset, sentence_id_to_idx)
    dependencies = minibatch_parse(sentences, model, eval_batch_size)

    UAS = all_tokens = 0.0
    with tqdm(total=len(dataset)) as prog:
        for i, ex in enumerate(dataset):
            head = [-1] * len(ex['word'])
            for h, t, in dependencies[i]:
                head[t] = h
            for pred_h, gold_h, gold_l, pos in \
                zip(head[1:], ex['head'][1:], ex['label'][1:], ex['pos'][1:]):
                assert self.id2tok[pos].startswith(P_PREFIX)
                pos_str = self.id2tok[pos][len(P_PREFIX):]
                if (self.with_punct) or (not punct(self.language, pos_str)):
                    UAS += 1 if pred_h == gold_h else 0
                    all_tokens += 1
            prog.update(i + 1)
    UAS /= all_tokens
    return UAS, dependencies

class ModelWrapper(object):

```



```

def __init__(self, parser, dataset, sentence_id_to_idx):
    self.parser = parser
    self.dataset = dataset
    self.sentence_id_to_idx = sentence_id_to_idx

def predict(self, partial_parses):
    mb_x = [self.parser.extract_features(p.stack, p.buffer, p.dependencies,
                                         self.dataset[self.sentence_id_to_idx[id]])
             for p in partial_parses]
    mb_x = np.array(mb_x).astype('int32')
    mb_x = torch.from_numpy(mb_x).long()
    mb_l = [self.parser.legal_labels(p.stack, p.buffer) for p in partial_parses]

    pred = self.parser.model(mb_x)
    pred = pred.detach().numpy()
    pred = np.argmax(pred + 10000 * np.array(mb_l).astype('float32'), 1)
    pred = ["S" if p == 2 else ("LA" if p == 0 else "RA") for p in pred]
    return pred

def read_conll(in_file, lowercase=False, max_example=None):
    examples = []
    with open(in_file) as f:
        word, pos, head, label = [], [], [], []
        for line in f.readlines():
            sp = line.strip().split('\t')
            if len(sp) == 10:
                if '-' not in sp[0]:
                    word.append(sp[1].lower() if lowercase else sp[1])
                    pos.append(sp[4])
                    head.append(int(sp[6]))
                    label.append(sp[7])
            elif len(word) > 0:
                examples.append({'word': word, 'pos': pos, 'head': head, 'label': label})
                word, pos, head, label = [], [], [], []
                if (max_example is not None) and (len(examples) == max_example):
                    break
        if len(word) > 0:
            examples.append({'word': word, 'pos': pos, 'head': head, 'label': label})
    return examples

def build_dict(keys, n_max=None, offset=0):
    count = Counter()
    for key in keys:
        count[key] += 1
    ls = count.most_common() if n_max is None \
        else count.most_common(n_max)

    return {w[0]: index + offset for (index, w) in enumerate(ls)}

def punct(language, pos):
    if language == 'english':
        return pos in ["'", ",", ".", ":", "`", "-LRB-", "-RRB-"]
    elif language == 'chinese':
        return pos == 'PU'
    elif language == 'french':
        return pos == 'PUNC'
    elif language == 'german':
        return pos in ["$.", "$,", "$[ "]

```

```

elif language == 'spanish':
    # http://nlp.stanford.edu/software/spanish-faq.shtml
    return pos in ["f0", "faa", "fat", "fc", "fd", "fe", "fg", "fh",
                  "fia", "fit", "fp", "fpa", "fpt", "fs", "ft",
                  "fx", "fz"]
elif language == 'universal':
    return pos == 'PUNCT'
else:
    raise ValueError('language: %s is not supported.' % language)

def minibatches(data, batch_size):
    x = np.array([d[0] for d in data])
    y = np.array([d[2] for d in data])
    one_hot = np.zeros((y.size, 3))
    one_hot[np.arange(y.size), y] = 1
    return get_minibatches([x, one_hot], batch_size)

def load_and_preprocess_data(reduced=True):
    config = Config()

    print("Loading data...",)
    start = time.time()
    train_set = read_conll(os.path.join(config.data_path, config.train_file),
                           lowercase=config.lowercase)
    dev_set = read_conll(os.path.join(config.data_path, config.dev_file),
                         lowercase=config.lowercase)
    test_set = read_conll(os.path.join(config.data_path, config.test_file),
                           lowercase=config.lowercase)

    if reduced:
        train_set = train_set[:1000]
        dev_set = dev_set[:500]
        test_set = test_set[:500]
    print("took {:.2f} seconds".format(time.time() - start))

    print("Building parser...",)
    start = time.time()
    parser = Parser(train_set)
    print("took {:.2f} seconds".format(time.time() - start))

    print("Loading pretrained embeddings...",)
    start = time.time()
    word_vectors = {}
    for line in open(config.embedding_file).readlines():
        sp = line.strip().split()
        word_vectors[sp[0]] = [float(x) for x in sp[1:]]
    embeddings_matrix = np.asarray(np.random.normal(0, 0.9, (parser.n_tokens, 50)),

    for token in parser.tok2id:
        i = parser.tok2id[token]
        if token in word_vectors:
            embeddings_matrix[i] = word_vectors[token]
        elif token.lower() in word_vectors:
            embeddings_matrix[i] = word_vectors[token.lower()]
    print("took {:.2f} seconds".format(time.time() - start))

    print("Vectorizing data...",)
    start = time.time()
    train_set = parser.vectorize(train_set)
    dev_set = parser.vectorize(dev_set)

```

```

test_set = parser.vectorize(test_set)
print("took {:.2f} seconds".format(time.time() - start))

print("Preprocessing training data...")
start = time.time()
train_examples = parser.create_instances(train_set)
print("took {:.2f} seconds".format(time.time() - start))

return parser, embeddings_matrix, train_examples, dev_set, test_set,

class AverageMeter(object):
    """Computes and stores the average and current value"""
    def __init__(self):
        self.reset()

    def reset(self):
        self.val = 0
        self.avg = 0
        self.sum = 0
        self.count = 0

    def update(self, val, n=1):
        self.val = val
        self.sum += val * n
        self.count += n
        self.avg = self.sum / self.count

if __name__ == '__main__':
    pass

```

So here is the skeleton code to implement this network using PyTorch. Complete the `__init__`, `embedding_lookup` and `forward` functions to implement the model.

Please DO NOT use `torch.nn.Linear` or `torch.nn.Embedding`. We are basically asking you to implement the Linear layer and Embedding layer by yourself so you can adjust the code according to the equation we got.

Please also follow the naming requirements in our TODO to avoid any problems.

## Hints

- Each of the variable (`self.embed_to_hidden_weight`, `self.embed_to_hidden_bias`, `self.hidden_to_logits_weight`, `self.hidden_to_logits_bias`) corresponds to ( $\mathbf{W}$ ,  $\mathbf{b}_1$ ,  $\mathbf{U}$ ,  $\mathbf{b}_2$ )
- It may help to work backwards in the algorithm (start from  $\hat{y}$ ) and keep track of the matrix/vector shapes
- At worst, loss should be smaller than 0.08, and UAS larger than 87 on the dev set (around :-)). The original paper got around 92.5 UAS.
- Should take around 1 hour to train the model on the entire dataset

In [6]:

```

#Linear class --> I created nn.Linear from scratch
#because torch.nn.Linear which is mentioned above,
#is not allowed to use

import math
import torch
from torch import Tensor
from torch.nn.parameter import Parameter, UninitializedParameter
from torch.nn import functional as F
from torch.nn import init
from torch.nn import Module

class Linear(Module):
    r"""Applies a linear transformation to the incoming data:  $y = xA^T + b$ 
    This module supports TensorFloat32 on ampere.
    Args:
        in_features: size of each input sample
        out_features: size of each output sample
        bias: If set to False, the layer will not learn an additive bias.
            Default: True
    Shape:
        - Input:  $(*, H_{in})$  where  $*$  means any number of
            dimensions including none and  $H_{in} = \text{in\_features}$ .
        - Output:  $(*, H_{out})$  where all but the last dimension
            are the same shape as the input and  $H_{out} = \text{out\_features}$ 
    Attributes:
        weight: the learnable weights of the module of shape
             $(\text{out\_features}, \text{in\_features})$ . The values are
            initialized from  $\mathcal{U}(-\sqrt{k}, \sqrt{k})$ , where
             $k = \frac{1}{\text{in\_features}}$ 
        bias: the learnable bias of the module of shape  $\text{out\_feature}$ 
            If attr: bias is True, the values are initialized from
             $\mathcal{U}(-\sqrt{k}, \sqrt{k})$  where
             $k = \frac{1}{\text{in\_features}}$ 
    Examples::
        >>> m = nn.Linear(20, 30)
        >>> input = torch.randn(128, 20)
        >>> output = m(input)
        >>> print(output.size())
        torch.Size([128, 30])
    """
    _constants_ = ['in_features', 'out_features']
    in_features: int
    out_features: int
    weight: Tensor

    def __init__(self, in_features: int, out_features: int, bias: bool = True,
                  device=None, dtype=None) -> None:
        factory_kwargs = {'device': device, 'dtype': dtype}
        super(Linear, self).__init__()
        self.in_features = in_features
        self.out_features = out_features
        self.weight = Parameter(torch.empty((out_features, in_features), **factory_k
        if bias:
            self.bias = Parameter(torch.empty(out_features, **factory_kwargs))
        else:
            self.register_parameter('bias', None)
        self.reset_parameters()

```

```
def reset_parameters(self) -> None:
    init.kaiming_uniform_(self.weight, a=math.sqrt(5))
    if self.bias is not None:
        fan_in, _ = init._calculate_fan_in_and_fan_out(self.weight)
        bound = 1 / math.sqrt(fan_in) if fan_in > 0 else 0
        init.uniform_(self.bias, -bound, bound)

def forward(self, input: Tensor) -> Tensor:
    return F.linear(input, self.weight, self.bias)

def extra_repr(self) -> str:
    return 'in_features={}, out_features={}, bias={}'.format(
        self.in_features, self.out_features, self.bias is not None
    )
```

In [7]:

```

class ParserModel(nn.Module):
    """ Feedforward neural network with an embedding layer and single hidden layer.
    The ParserModel will predict which transition should be applied to a
    given partial parse configuration.
    PyTorch Notes:
        - Note that "ParserModel" is a subclass of the "nn.Module" class. In
          PyTorch all neural networks
            are a subclass of this "nn.Module".
        - The "__init__" method is where you define all the layers and their
          respective parameters (embedding layers, linear layers, dropout layers, et
        - "__init__" gets automatically called when you create a new instance
          of your class, e.g. when you write "m = ParserModel()".
        - Other methods of ParserModel can access variables that have "self."
          prefix. Thus, you should add the "self." prefix layers, values, etc.
          that you want to utilize in other ParserModel methods.
        - For further documentation on "nn.Module" please see
          https://pytorch.org/docs/stable/nn.html.
    """

    def __init__(self, embeddings, n_features=36,
                  hidden_size=400, n_classes=3, dropout_prob=0.5):
        """ Initialize the parser model.
        @param embeddings (Tensor): word embeddings (num_words, embedding_size)
        @param n_features (int): number of input features
        @param hidden_size (int): number of hidden units
        @param n_classes (int): number of output classes
        @param dropout_prob (float): dropout probability
        """

        super(ParserModel, self).__init__()
        self.n_features = n_features
        self.n_classes = n_classes
        self.dropout_prob = dropout_prob
        self.embed_size = embeddings.shape[1]
        self.hidden_size = hidden_size
        self.pretrained_embeddings = nn.Embedding(
            embeddings.shape[0], self.embed_size)
        self.pretrained_embeddings.weight = nn.Parameter(
            torch.tensor(embeddings))

        # YOUR CODE HERE (~5 Lines)
        # TODO:
        # 1) Construct `self.embed_to_hidden` linear layer, initializing the weight
        # with the `nn.init.xavier_uniform` function with `gain = 1` (default)
        # 2) Construct `self.dropout` layer.
        # 3) Construct `self.hidden_to_logits` linear layer, initializing the weight
        # with the `nn.init.xavier_uniform` function with `gain = 1` (default)
        ###
        # Note: Here, we use Xavier Uniform Initialization for our Weight initialization
        # It has been shown empirically, that this provides better initial weights
        # for training networks than random uniform initialization.
        # For more details checkout this great blogpost:
        # http://andyljones.tumblr.com/post/110998971763/an-explanation-of-xavier-in
        # Hints:
        # - After you create a linear layer you can access the weight
        # matrix via:
        # linear_layer.weight
        ###
        # Please see the following docs for support:
        # Linear Layer: https://pytorch.org/docs/stable/nn.html#torch.nn.Linear

```

```
# Xavier Init: https://pytorch.org/docs/stable/nn.html#torch.nn.init.xavier\_uniform
# Dropout: https://pytorch.org/docs/stable/nn.html#torch.nn.Dropout
```

```
# self.embed_to_hidden = nn.Linear(self.n_features * self.embed_size, self.hidden_size)
# nn.init.xavier_uniform_(self.embed_to_hidden.weight, gain=1)
# self.dropout = nn.Dropout(self.dropout_prob)
# self.hidden_to_logits = nn.Linear(self.hidden_size, self.n_classes)
# nn.init.xavier_uniform_(self.hidden_to_logits.weight, gain=1)
```

```
# calling Linear class from the above
self.embed_to_hidden = Linear(self.n_features * self.embed_size, self.hidden_size)
nn.init.xavier_uniform_(self.embed_to_hidden.weight, gain=1)
self.dropout = nn.Dropout(self.dropout_prob)
self.hidden_to_logits = Linear(self.hidden_size, self.n_classes)
nn.init.xavier_uniform_(self.hidden_to_logits.weight, gain=1)
```

```
# END YOUR CODE
```

```
def embedding_lookup(self, t):
    """ Utilize `self.pretrained_embeddings` to map input `t` from input tokens
        (integers) to embedding vectors.
        PyTorch Notes:
            - `self.pretrained_embeddings` is a torch.nn.Embedding object
              that we defined in __init__
            - Here `t` is a tensor where each row represents a list of
              features. Each feature is represented by an integer (input token).
            - In PyTorch the Embedding object, e.g.
              `self.pretrained_embeddings`, allows you to
              go from an index to embedding. Please see the documentation
              (https://pytorch.org/docs/stable/nn.html#torch.nn.Embedding)
              to learn how to use `self.pretrained_embeddings` to extract
              the embeddings for your tensor `t`.
        @param t (Tensor): input tensor of tokens (batch_size, n_features)
        @return x (Tensor): tensor of embeddings for words represented in t
                           (batch_size, n_features * embed_size)
    """
    # YOUR CODE HERE (~1-3 Lines)
    # TODO:
    # 1) Use `self.pretrained_embeddings` to lookup the embeddings for the
    # input tokens in `t`.
    # 2) After you apply the embedding lookup, you will have a tensor shape
    # (batch_size, n_features, embedding_size).
    # Use the tensor `view` or `reshape` method to reshape the embeddings tensor
    # (batch_size, n_features * embedding_size)
    ###
    # Note: In order to get batch_size, you may need use the tensor .size()
    # function:
    # https://pytorch.org/docs/stable/tensors.html#torch.Tensor.size
    ###
    # Please see the following docs for support:
    # Embedding Layer: https://pytorch.org/docs/stable/nn.html#torch.nn.Embedding
    # View: https://pytorch.org/docs/stable/tensors.html#torch.Tensor.view

    embeddings = self.pretrained_embeddings.forward(t)
    x = embeddings.view(t.shape[0], self.embed_size * self.n_features)
```

```

# END YOUR CODE
return x

def forward(self, t):
    """ Run the model forward.
    Note that we will not apply the softmax function here because it is
    included in the loss function nn.CrossEntropyLoss
    PyTorch Notes:
        - Every nn.Module object (PyTorch model) has a `forward`
          function.
        - When you apply your nn.Module to an input tensor `t` this
          function is applied to the tensor.
          For example, if you created an instance of your ParserModel
          and applied it to some `t` as follows, the `forward` function
          would called on `t` and the result would be stored in the
          `output` variable:
              model = ParserModel()
              output = model(t) # this calls the forward function
    # torch.nn.Module.forward
        - For more details checkout: https://pytorch.org/docs/stable/nn.html
    @param t (Tensor): input tensor of tokens (batch_size, n_features)
    @return logits (Tensor): tensor of predictions (output after applying
                             the layers of the network) without applying
                             softmax (batch_size, n_classes)

    """
    # YOUR CODE HERE (~3-5 lines)
    # TODO:
    # 1) Apply `self.embedding_lookup` to `t` to get the embeddings
    # 2) Apply `embed_to_hidden` linear layer to the embeddings
    # 3) Apply relu non-linearity to the output of step 2 to get the hidden unit
    # 4) Apply dropout layer to the output of step 3.
    # 5) Apply `hidden_to_logits` layer to the output of step 4 to get the logits
    ###
    # Note: We do not apply the softmax to the logits here, because
    # the loss function (torch.nn.CrossEntropyLoss) applies it more efficiently.
    ###
    # Please see the following docs for support:
    # ReLU: https://pytorch.org/docs/stable/nn.html?highlight=relu#torch.nn.func

    embed = self.embedding_lookup(t)
    hidden = F.relu(self.embed_to_hidden(embed))
    dropout = self.dropout(hidden)
    logits = self.hidden_to_logits(dropout)

    # END YOUR CODE
    return logits

```

Now complete the `train_for_epoch` and `train` functions to actually train the model.



In [8]:

```

# -----
# Primary Functions
# -----

def train(parser, train_data, dev_data, output_path, batch_size=1024, n_epochs=10, lr=0.001):
    """ Train the neural dependency parser.
    @param parser (Parser): Neural Dependency Parser
    @param train_data ():
    @param dev_data ():
    @param output_path (str): Path to which model weights and results are written.
    @param batch_size (int): Number of examples in a single batch
    @param n_epochs (int): Number of training epochs
    @param lr (float): Learning rate
    """
    best_dev_UAS = 0

    # YOUR CODE HERE (2 lines)
    # TODO:
    # 1) Construct Adam Optimizer in variable `optimizer`
    # 2) Construct the Cross Entropy Loss Function in variable `loss_func`
    ###
    # Hint: Use `parser.model.parameters()` to pass optimizer
    # necessary parameters to tune.
    # Please see the following docs for support:
    # Adam Optimizer: https://pytorch.org/docs/stable/optim.html
    # Cross Entropy Loss: https://pytorch.org/docs/stable/nn.html#crossentropyloss

    optimizer = optim.Adam(parser.model.parameters(), lr=lr)
    loss_func = nn.CrossEntropyLoss()

    # END YOUR CODE

    for epoch in range(n_epochs):
        print("Epoch {}: out of {}".format(epoch + 1, n_epochs))
        dev_UAS = train_for_epoch(
            parser, train_data, dev_data, optimizer, loss_func, batch_size)
        if dev_UAS > best_dev_UAS:
            best_dev_UAS = dev_UAS
            print("New best dev UAS! Saving model.")
            torch.save(parser.model.state_dict(), output_path)
        print("")

def train_for_epoch(parser, train_data, dev_data, optimizer, loss_func, batch_size):
    """ Train the neural dependency parser for single epoch.
    Note: In PyTorch we can signify train versus test and automatically have
    the Dropout Layer applied and removed, accordingly, by specifying
    whether we are training, `model.train()`, or evaluating, `model.eval()`
    @param parser (Parser): Neural Dependency Parser
    @param train_data ():
    @param dev_data ():
    @param optimizer (nn.Optimizer): Adam Optimizer
    @param loss_func (nn.CrossEntropyLoss): Cross Entropy Loss Function

```

```

@param batch_size (int): batch size
@param lr (float): learning rate
@return dev_UAS (float): Unlabeled Attachment Score (UAS) for dev data
"""

parser.model.train() # Places model in "train" mode, i.e. apply dropout layer
n_minibatches = math.ceil(len(train_data) / batch_size)
loss_meter = AverageMeter()

with tqdm(total=(n_minibatches)) as prog:
    for i, (train_x, train_y) in enumerate(minibatches(train_data, batch_size)):
        optimizer.zero_grad() # remove any baggage in the optimizer
        loss = 0. # store loss for this batch here
        train_x = torch.from_numpy(train_x).long()
        train_y = torch.from_numpy(train_y.nonzero()[1]).long()

        # YOUR CODE HERE (~5-10 lines)
        # TODO:
        # 1) Run train_x forward through model to produce `logits`
        # 2) Use the `loss_func` parameter to apply the PyTorch CrossEntropyLoss
        # This will take `logits` and `train_y` as inputs. It will output the CrossEntropyLoss
        # between softmax(`logits`) and `train_y`. Remember that softmax(`logits`)
        # are the predictions (y^ from the PDF).
        # 3) Backprop losses
        # 4) Take step with the optimizer
        # Please see the following docs for support:
        # Optimizer Step: https://pytorch.org/docs/stable/optim.html#optimizer-s

        logits = parser.model.forward(train_x)
        loss = loss_func(logits, train_y)
        loss.backward()
        optimizer.step()

        # END YOUR CODE
        prog.update(1)
        loss_meter.update(loss.item())

print("Average Train Loss: {}".format(loss_meter.avg))

print("Evaluating on dev set",)
parser.model.eval() # Places model in "eval" mode, i.e. don't apply dropout layer
dev_UAS, _ = parser.parse(dev_data)
print("- dev UAS: {:.2f}".format(dev_UAS * 100.0))
return dev_UAS

```

## Part 5: Evaluation

Now execute this code to actually train your model and compute predictions on test data from Penn Treebank (annotated with Universal Dependencies)

In [9]:

```

# Note: Set debug to False, when training on entire corpus
# debug = True
debug = False

print(80 * "=")
print("INITIALIZING")
print(80 * "=")
parser, embeddings, train_data, dev_data, test_data = load_and_preprocess_data(
    debug)

start = time.time()
model = ParserModel(embeddings)
parser.model = model
print("took {:.2f} seconds\n".format(time.time() - start))

print(80 * "=")
print("TRAINING")
print(80 * "=")
output_dir = "results/{:%Y%m%d_%H%M%S}/".format(datetime.now())
output_path = output_dir + "model.weights"

if not os.path.exists(output_dir):
    os.makedirs(output_dir)

train(parser, train_data, dev_data, output_path,
      batch_size=1024, n_epochs=10, lr=0.0005)

if not debug:
    print(80 * "=")
    print("TESTING")
    print(80 * "=")
    print("Restoring the best model weights found on the dev set")
    parser.model.load_state_dict(torch.load(output_path))
    print("Final evaluation on test set",)
    parser.model.eval()
    UAS, dependencies = parser.parse(test_data)
    print("- test UAS: {:.2f}".format(UAS * 100.0))
    print("Done!")

```

```

=====
=====
INITIALIZING
=====
=====
Loading data...
took 0.93 seconds
Building parser...
took 0.52 seconds
Loading pretrained embeddings...
took 1.23 seconds
Vectorizing data...
took 0.72 seconds
Preprocessing training data...
took 20.94 seconds
took 0.04 seconds

=====
=====
TRAINING

```





```
=====
```

```
TESTING
```

```
=====
```

```
=====
```

```
Restoring the best model weights found on the dev set
```

```
Final evaluation on test set
```

```
2919736it [00:00, 133773121.24it/s]
```

```
- test UAS: 89.82
```

```
Done!
```

```
In [ ]:
```

```
In [ ]:
```